

# Chapter 5: Splines

## 5.1 Smoothing Splines

### Natural Cubic Splines

If we order our points in increasing value a natural cubic spline is a function made up of sections of cubic polynomials linking each successive pair of points. Of all the functions that are continuous and interpolate the points the “smoothest” minimizes  $\int_{x_1}^{x_n} f''(x)^2 dx$ .

### Cubic Smoothing Splines

Because most data is noisy and you don't have a full grasp of the data-generating mechanism you usually want to smooth the data rather than interpolate the points. So instead of fixing  $g(x_i) = y_i$  we set them as parameters and attempt to minimize

$$\sum_{i=1}^n (y_i - g(x_i))^2 + \lambda \int g''(x) dx$$

with  $\lambda$  as the tuning parameter. We call  $g(x)$  a *smoothing spline*.

While these are ideal smoothers (smoothing function that minimizes error) they have as many free parameters as there are data so can be problematic to estimate. Also because we are smoothing the function anyway we almost certainly won't need all  $n$  parameters.

## 5.2 Penalized Regression Splines

A Regression spline constructs a spline basis and then uses that basis to model the original data set.

## 5.3 Some One-Dimensional Smoothers

### Cubic Splines

Already talked about these. If I get adventurous I'll try to fit one manually

From `mgcv` these can be called with `s(x, bs = "cr")`

### Cyclic splines

These match the outermost knots so that the smooth function is equal there. Think of modeling smooth effects throughout the year, you wouldn't want there to be a discontinuity at the end of the year.

### B-splines

A B-Spline (Basis Spline) are only non-zero between  $m + 3$  adjacent knots ( $m + 1$  is the order of the basis, so for cubic splines  $m = 2$  and a B-spline would be non-zero for the nearest 5 knots). This makes them stable to compute. We can define them recursively by saying any  $m$  order spline is a weighted sum of lower order splines where the weights are defined by how close the  $x$ -values are to the series of knots. Wikipedia says:

(the first) ramps from zero to one as  $x$  goes from  $t_{\{i\}}$  to  $t_{\{i+k\}}$  and (the second) ramps from one to zero as  $x$  goes from  $t_{\{i+1\}}$  to  $t_{\{i+k+1\}}$ .

The final stage of recursion is just a binary indicating which knot span  $x$  is in.

So if we want a cubic B-spline with 6 knots we actually need to define 10 knots; 6 for the locations, 2 for the ends, and 2 for the cubic order.

```
library(ggplot2)
library(dplyr)
```

```
##
## Attaching package: 'dplyr'

## The following objects are masked from 'package:stats':
##
##   filter, lag

## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union
```

```
library(tidyr)
x <- 1:100 + rnorm(100)
k <- seq(0, 100, length.out = 10)
```

```
print(x)
```

```
##   [1]  2.8304466  0.6383422  3.4615389  6.1504733  3.4152325
##   [6]  6.2263610  6.5477073  9.0993343  8.0418226  9.3717109
##  [11] 10.9756387 10.1283545 12.2468605 12.6780563 16.3155565
##  [16] 15.5912416 16.2977648 16.0640794 18.9211216 18.7762839
##  [21] 21.5907149 23.2462060 23.4015386 24.1061468 27.0100930
##  [26] 26.2023187 27.0883877 26.6310496 30.2142395 30.8753991
##  [31] 32.0275432 31.2226954 33.7821897 33.0819412 34.6630579
##  [36] 37.3655724 36.1177560 38.4240626 39.7277989 39.6909017
##  [41] 41.9016775 40.8365454 43.6544693 45.1888068 45.0414979
##  [46] 45.8185781 47.4888263 48.8236175 48.8453556 50.0763945
##  [51] 51.0970741 51.8023765 51.5922700 54.7125974 55.7241879
##  [56] 54.7663641 55.4921620 56.6276969 60.3829692 58.2453033
##  [61] 60.5776437 61.7128858 62.2392796 64.9088296 67.8037814
##  [66] 66.1063845 67.0834966 68.8881896 69.6538935 69.9908926
##  [71] 72.0177223 71.9256885 72.9552996 74.8237473 73.5271670
##  [76] 77.2170905 78.0646884 77.2505063 79.1991229 80.2406643
##  [81] 80.8950586 82.7432086 82.3374130 85.5897813 82.5047999
##  [86] 85.7040926 87.4424702 86.8896348 89.4160749 89.9235826
##  [91] 90.2057115 90.8331037 92.0501123 92.3223260 95.6368462
##  [96] 96.4862053 97.3810915 99.1371482 100.5826330 101.4176250
```

```
print(k)
```

```
##   [1]  0.00000 11.11111 22.22222 33.33333 44.44444 55.55556 66.66667
##   [8] 77.77778 88.88889 100.00000
```

```
bspline <- function(x, k, i, m = 2){
  ## Evaluate ith B-spline basis function of order m at the
  ## values in x, given knot locations in k
  if(m == -1){
    res <- as.numeric(x < k[i + 1] & x >= k[i])
```

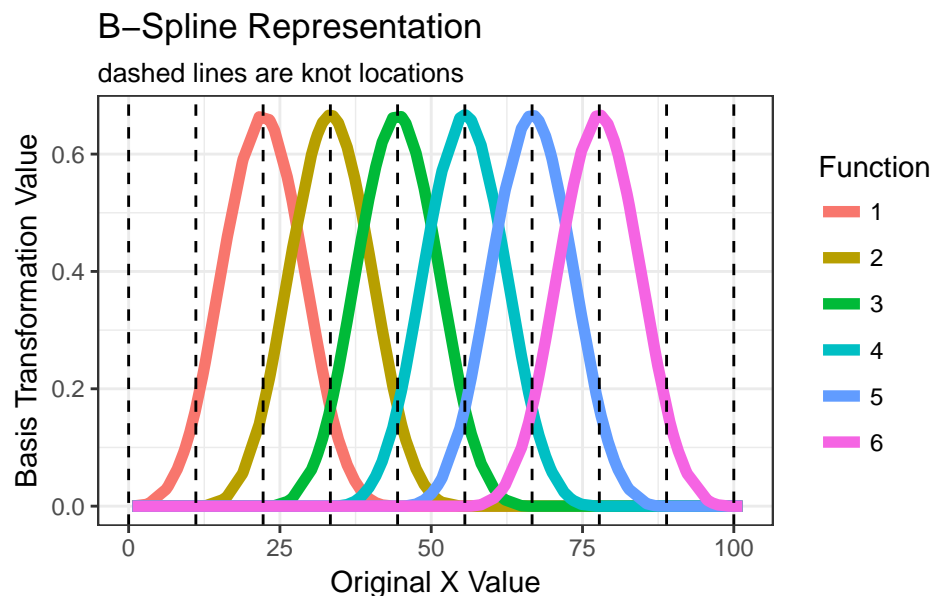
```

} else {
  z0 <- (x - k[i]) / (k[i + m + 2] - k[i + 1])
  z1 <- (k[i + m + 2] - x) / (k[i + m + 2] - k[i + 1])
  res <- z0 * bspline(x, k, i, m - 1) + z1 * bspline(x, k, i + 1, m - 1)
}
return(res)
}

spline_basis <- vapply(1:6, function(i, x_ = x, k_ = k) bspline(x = x_, k = k_, i = i), numeric(100))
spline_basis_df <- data.frame(spline_basis, stringsAsFactors = F)
names(spline_basis_df) <- as.character(seq_len(ncol(spline_basis)))
spline_basis_df$X <- x

spline_basis_df %>%
  gather(Function, Value, -X) %>%
  ggplot(aes(x = X, y = Value, color = Function)) +
  geom_line(size = 2) +
  geom_vline(aes(xintercept = loc), data = data_frame(loc = k), linetype = "dashed") +
  ggtitle("B-Spline Representation", "dashed lines are knot locations") +
  theme_bw() +
  xlab("Original X Value") +
  ylab("Basis Transformation Value")

```



So as you can see the outside 2 knots on both sides don't actually have full coverage, which is why we need to put knot locations outside the range of our data. We would then use this new basis as the variables in our regression model. We would then use these transformed values as inputs for our regression model.

```
knitr::kable(head(spline_basis_df), align = "c")
```

	1	2	3	4	5	6	X
0.0027551	0	0	0	0	0	0	2.8304466
0.0000316	0	0	0	0	0	0	0.6383422
0.0050395	0	0	0	0	0	0	3.4615389
0.0282684	0	0	0	0	0	0	6.1504733

	1	2	3	4	5	6	X
	0.0048399	0	0	0	0	0	3.4152325
	0.0293278	0	0	0	0	0	6.2263610

## P-Splines

P-splines are low rank smoothers using a B-spline basis, but with a *difference penalty* applied to the parameters to control wiggleness. This means we are penalizing the squared differences between adjacent  $\beta_i$  values. We can represent this penalty as

$$\sum_{i=1}^{k-1} (\beta_{i+1} - \beta_i)^2 = \beta^T \mathbf{P}^T \mathbf{P} \beta$$

where  $\mathbf{P}$  is just a diagonal difference matrix.

```
k <- 6
P <- diff(diag(k), differences = 1)
S <- t(P) %*% P
S
```

```
##      [,1] [,2] [,3] [,4] [,5] [,6]
## [1,]    1   -1    0    0    0    0
## [2,]   -1    2   -1    0    0    0
## [3,]    0   -1    2   -1    0    0
## [4,]    0    0   -1    2   -1    0
## [5,]    0    0    0   -1    2   -1
## [6,]    0    0    0    0   -1    1
```

P-Splines do require evenly spaced knots, but other than that they are very flexible.

From `mgcv` these can be called with `s(x, bs = "ps", m = c(2, 3))` where  $m$  is the order for the basis and penalties, respectively.

## Adaptive Smoothers

Sometimes we want the amount of smoothing to vary along with the  $x$  value. We can do this by adding weights to the differences penalties and letting those weights vary smoothly with  $x$ .

From `mgcv` these can be called with `s(x, bs = "ad", k = 40, m = 4)`.

## SCOP Splines

Can add shape constraints, aka monotonic functions.

## 5.4 Some Useful Smoother Theory

### Identifiability Constraints

Since each smooth can't have it's own intercept so we force the smooth terms to sum to 0 over the observed values of  $x$ .

## Effective Degrees of Freedom

The effective degrees of freedom from a smooth is

$$\sum_i (1 + \lambda D_{ii})^{-1}$$

where  $\lambda$  is the smoothing parameter and  $D$  is a diagonal matrix of eigen values of the  $\mathbf{R}^{-\mathbf{T}}\mathbf{S}\mathbf{R}^{-1}$ . We can rewrite the degrees of freedom as the trace of  $\mathbf{F}$  where

$$\mathbf{F} = (\mathbf{X}^{\mathbf{T}}\mathbf{X} + \lambda\mathbf{S})^{-1}\mathbf{X}^{\mathbf{T}}\mathbf{X}$$

So the degrees of freedom when there is no smoothing ( $\lambda = 0$ ) is just the number of coefficients in the model and the number of zero eigenvalues of the penalty ( $\lambda \rightarrow \infty$ ).

## Null Space Penalties

Most smoothing penalties treat some null space of functions as completely smooth and therefore have zero penalty; A Cubic spline penalty ( $\int f''(x)^2 dx$ ) is zero for any straight line (2nd derivate is 0, so there is nothing to sum to penalize). So when the penalty approaches infinity the smoother does not tend to 0 effect but tends to a straight line! So the penalty is not enough to remove a smooth term from the model altogether.

We can alleviate this by adding an extra penalty which only penalizes functions in the penalty null space (where a smoothing penalty has no effect).

The `select` argument in `gam` can be used to apply such penalties:

If this is TRUE then gam can add an extra penalty to each term so that it can be penalized to zero. This means that the smoothing parameter estimation that is part of fitting can completely remove terms from the model. If the corresponding smoothing parameter is estimated as zero then the extra penalty has no effect.

## 5.5 Isotropic Smoothing

Isotropic smooths will produce identical predictions of the response variable under any rotations or reflections of covariates.

### Thin Plate Regression Splines

So far each smoothing basis has the following characteristics:

1. You have to choose knot locations
2. Each basis can only incorporate one variable
3. It is not clear that they are better than any other basis

### Thin Plate Splines

suppose we have a smooth function  $g(x)$  we would like to estimate from  $n$  observations where  $y_i = g(x_i) + \epsilon_i$ . Thin plate splines estimate  $g$  by finding the function  $f$  that minimizes

$$\|\mathbf{y} - \mathbf{f}\|^2 + \lambda J_{md}(f)$$

where  $J_{md}(f)$  is a penalty function measuring the “wiggleness” of  $f$ .

The functions making up the function space,  $\mathbf{f}$ , are linear independent polynomials spanning the space of the polynomials in  $\mathbb{R}^d$ . Also the first couple functions (depending on the exact rank and dimension) span the space of functions for which  $J_{md}(f)$  is 0, i.e., are in the null space of  $J_{md}(f)$ . For example if  $m = d = 2$  then  $\phi_1(x) = 1$ ,  $\phi_2(x) = x_1$ , and  $\phi_3(x) = x_2$ .

We do not have to define knots or select the basis functions for thin plate splines. Also we can use as many predictors as we like.

The problem is that these are very computationally costly; there are as many unknown parameters as data. So we want a low rank approximation of these splines

### Thin Plate Regression Splines

Basically we want to truncate some of the wiggly components of the thin plate spline.

## 5.6 Tensor Product Smooth Interactions

Tensor products are scale invariant. A thin plate spline works similar to a flexible strip of and scales up to a flexible sheet. Tensor products just interlock multiple strips.

### Tensor Product Bases

Tensor products, basically, set a smooth function of one covariate using a sequence of knots. Then we let each parameter of that smooth vary with  $z$  as well by defining them as a smooth function of  $z$ . The same tensor product would be found if we started with  $z$  instead of  $x$ .

These are scale invariate.

From `mgcv` these can be called with the `te` function, `te(x, y, z)`

### ANOVA Decompositions of Smooths

Suppose we want to test a smooth interaction

$$f_1(x) + f_2(z) + f_3(x, z)$$

we can build tensor product interaction smooths with sum-to-zero constraints so that any main effects are removed.

In `mgcv` these can be called with the `ti` function.

## 5.7 Isotropy Versus Scale Invariance

Isotropic smooths are sensitive to linear rescaling of a single covariate. The reason is that the thin plate spline attempts to achieve the same smoothness per unit change in every covariate, so when a unit changes scale the function starts to fall apart. Tensor product smooths are not affected by any re-scaling.